Assignment-11

Problem Definition: Build scaling plans in AWS that balance the load on different EC2 instance.

Step-1: Signing in to the AWS Management console,go to lauch template.A template name should be given along with a description and the auto-scaling option is to be checked.Then ubuntu is selected from quickstart.

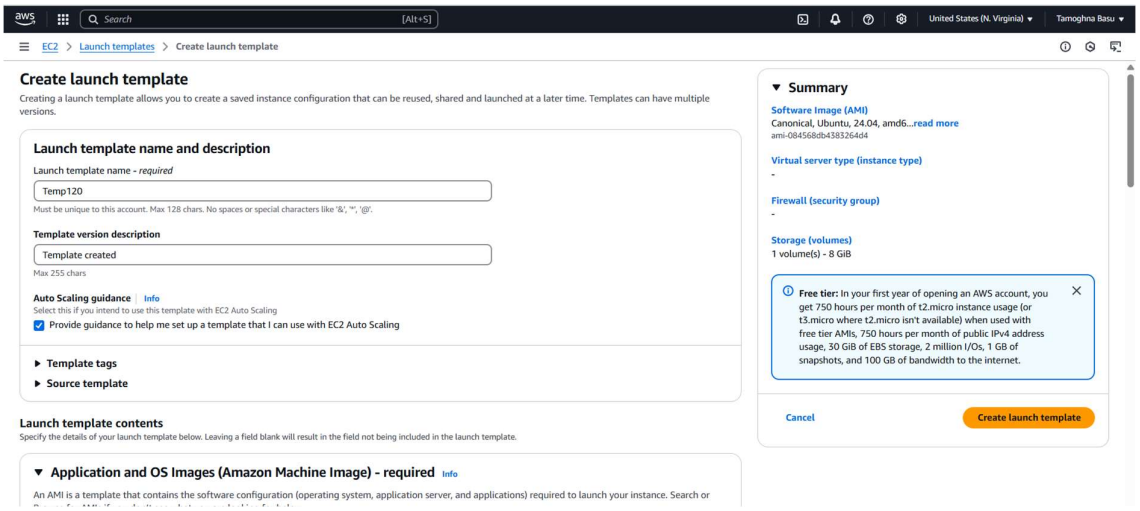

Step-2: Instance type is selected as t2.micro.

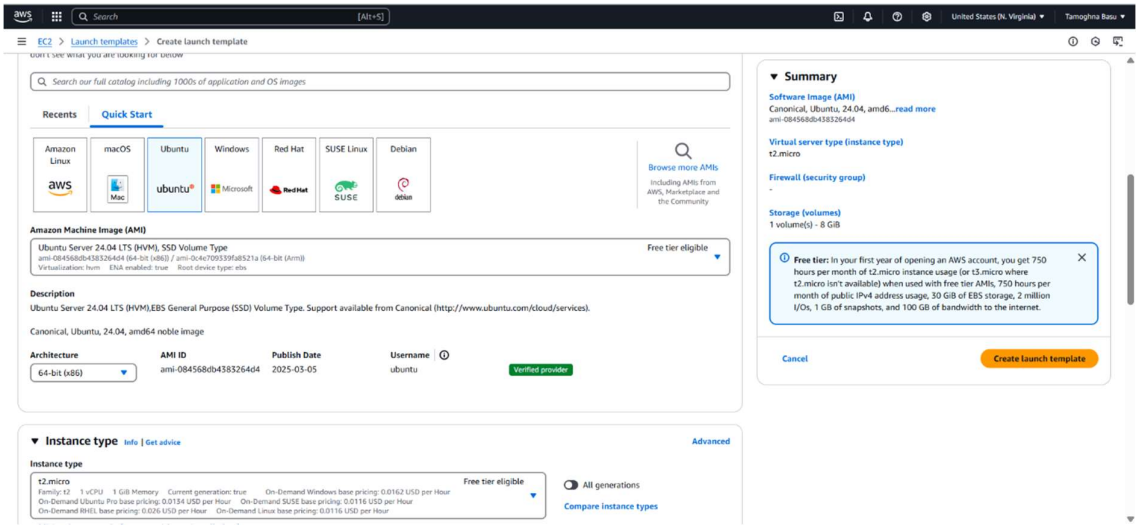

Step-3:A key-pair is selected along with the secuirity group and the user data is given.Create launch template is selected to create a new template.

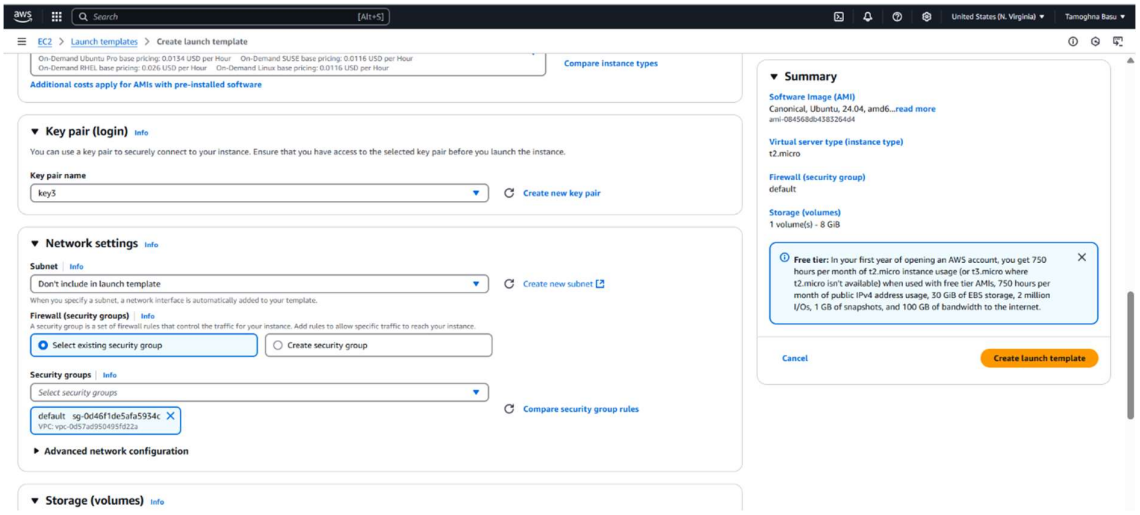

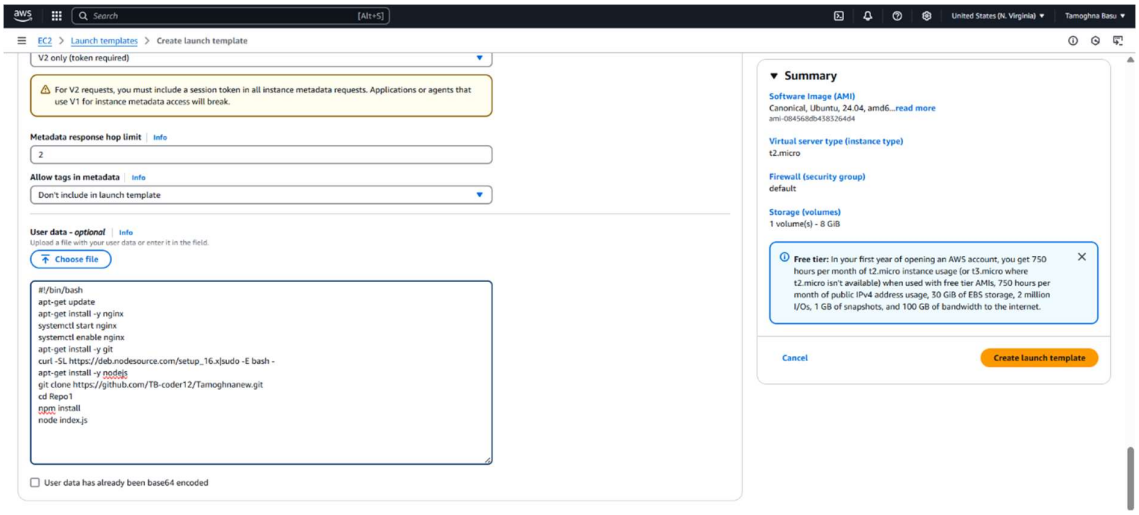

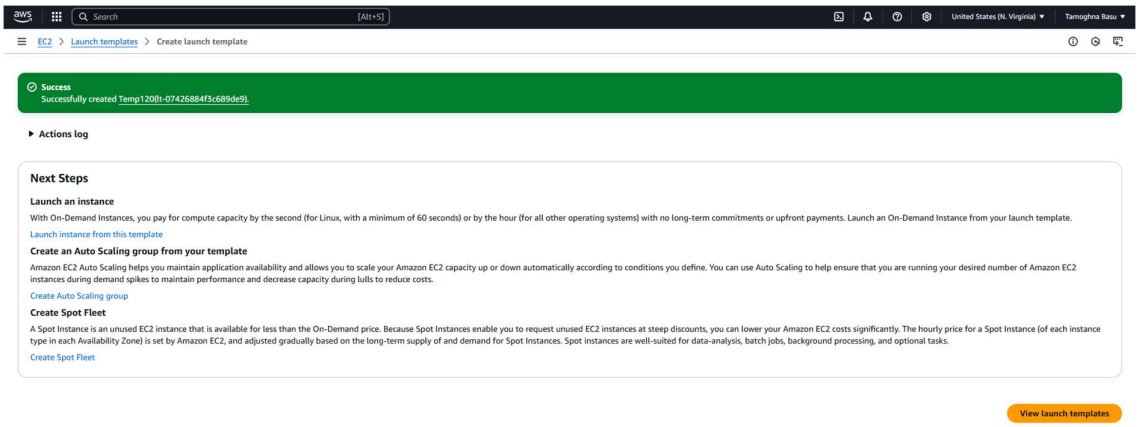

Step-4:Auto-scaling group is reached ,a name is given,template is selected,zones are selected,along with Balanced best effort.

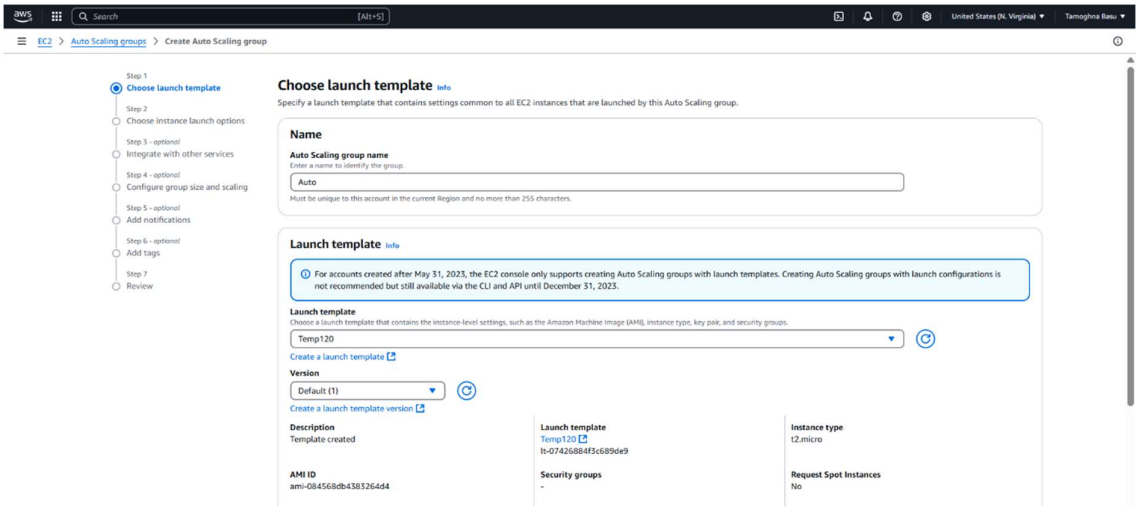

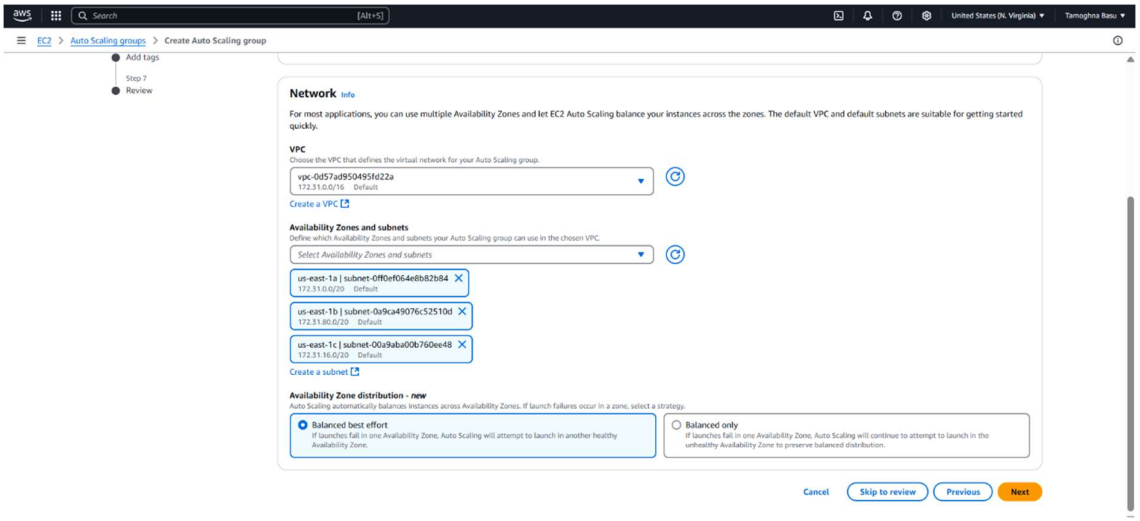

Step-5:On the next page,attach to a new load balancer is selected along with Internet-facing as the load balancer scheme.Additional health check types is enabled.

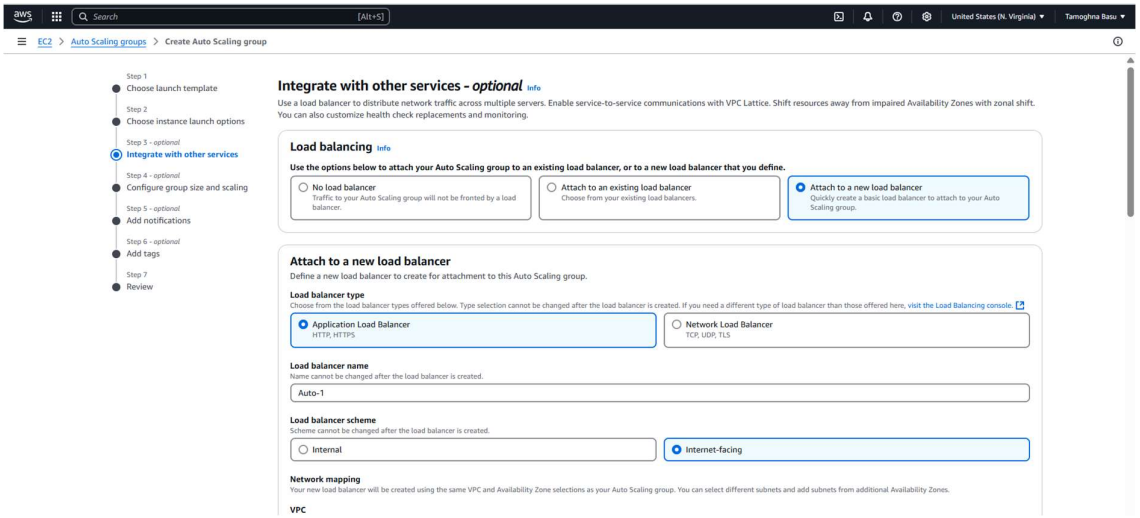
Search [Alt+S]
United States (N. Virginia)
Tamoghna Basu
EC2 > Auto Scaling groups > Create Auto Scaling group
Step 1
Choose launch template
Step 2
Choose instance launch options
Step 3 - optional
Integrate with other services
Step 4 - optional
Configure group size and scaling
Step 5 - optional
Add notifications
Step 6 - optional
Add tags
Step 7
Review
Integrate with other services - optional Info
Use a load balancer to distribute network traffic across multiple servers. Enable service-to-service communications with VPC Lattice. Shift resources away from impaired Availability Zones with zonal shift. You can also customize health check replacements and monitoring.
Load balancing Info
Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.
No load balancer
Traffic to your Auto Scaling group will not be fronted by a load balancer.
Attach to an existing load balancer
Choose from your existing load balancers.
Attach to a new load balancer
Quickly create a basic load balancer to attach to your Auto Scaling group.
Attach to a new load balancer
Define a new load balancer to create for attachment to this Auto Scaling group.
Load balancer type
Choose from the load balancer types offered below. Type selection cannot be changed after the load balancer is created. If you need a different type of load balancer than those offered here, visit the Load Balancing console.
Application Load Balancer
HTTP, HTTPS
Network Load Balancer
TCP, UDP, TLS
Load balancer name
Name cannot be changed after the load balancer is created.
Auto-1
Load balancer scheme
Scheme cannot be changed after the load balancer is created.
Internal
Internet-facing
Network mapping
Your new load balancer will be created using the same VPC and Availability Zone selections as your Auto Scaling group. You can select different subnets and add subnets from additional Availability Zones.
VPC

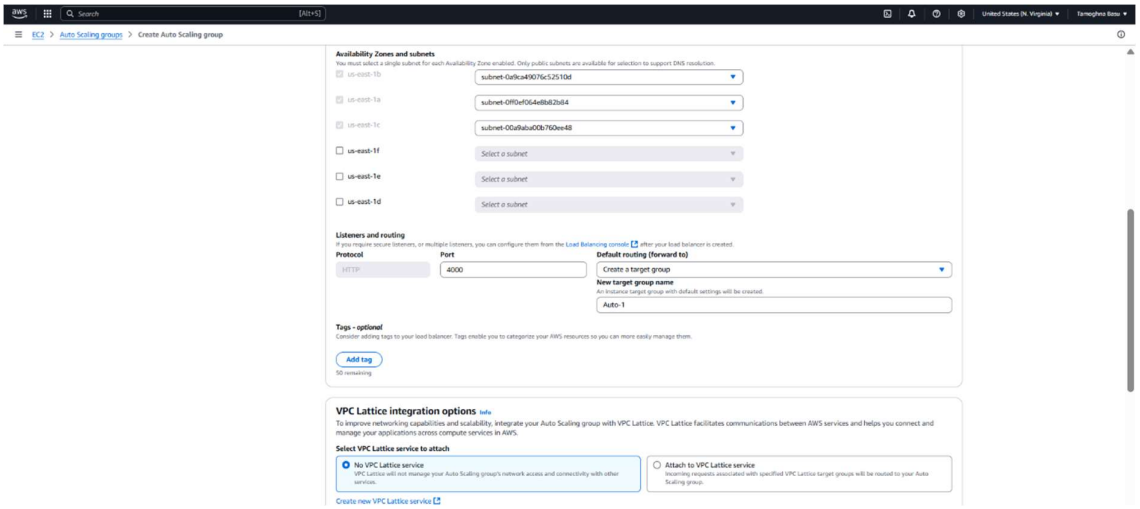
Search [Alt+S]
United States (N. Virginia)
Tamoghna Basu
EC2 > Auto Scaling groups > Create Auto Scaling group
Availability Zones and subnets
You must select a single subnet for each Availability Zone enabled. Only public subnets are available for selection to support DNS resolution.
us-east-1b
subnet-0a9ca49076e52510d
us-east-1a
subnet-0ff0ef064e8b82b84
us-east-1c
subnet-00a9aba00b760ee48
us-east-1f
Select a subnet
us-east-1e
Select a subnet
us-east-1d
Select a subnet
Listeners and routing
If you require secure listeners, or multiple listeners, you can configure them from the Load Balancing console after your load balancer is created.
Protocol
Port
Default routing (forward to)
HTTP
4000
Create a target group
New target group name
An instance target group with default settings will be created.
Auto-1
Tags - optional
Consider adding tags to your load balancer. Tags enable you to categorize your AWS resources so you can more easily manage them.
Add tag
50 remaining
VPC Lattice integration options Info
To improve networking capabilities and scalability, integrate your Auto Scaling group with VPC Lattice. VPC Lattice facilitates communications between AWS services and helps you connect and manage your applications across compute services in AWS.
Select VPC Lattice service to attach
No VPC Lattice service
VPC Lattice will not manage your Auto Scaling group's network access and connectivity with other services.
Attach to VPC Lattice service
Incoming requests associated with specified VPC Lattice target groups will be routed to your Auto Scaling group.
Create new VPC Lattice service

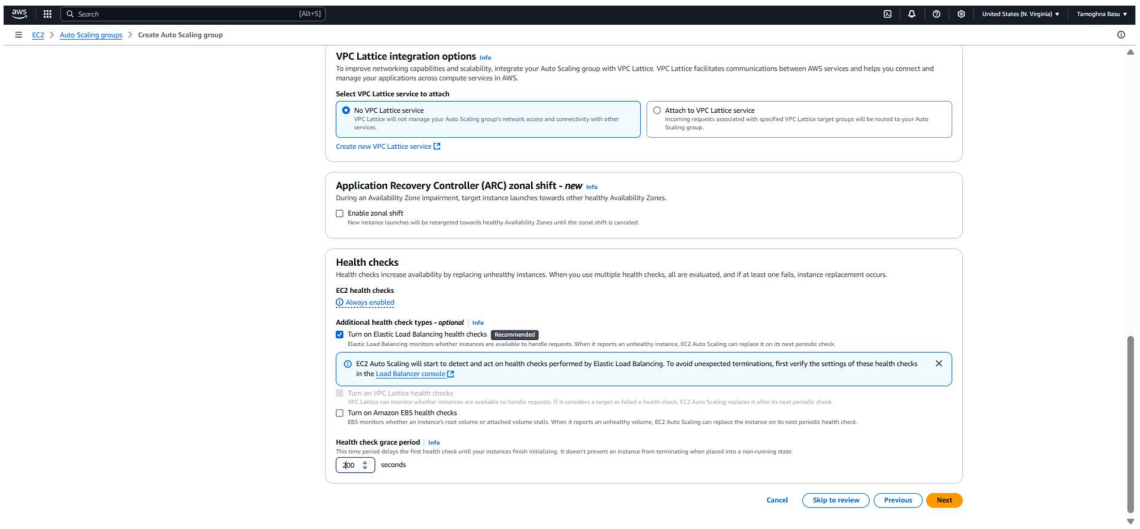
Search [Alt+S]
United States (N. Virginia)
Tamoghna Basu
EC2 > Auto Scaling groups > Create Auto Scaling group
VPC Lattice integration options Info
To improve networking capabilities and scalability, integrate your Auto Scaling group with VPC Lattice. VPC Lattice facilitates communications between AWS services and helps you connect and manage your applications across compute services in AWS.
Select VPC Lattice service to attach
No VPC Lattice service
VPC Lattice will not manage your Auto Scaling group's network access and connectivity with other services.
Attach to VPC Lattice service
Incoming requests associated with specified VPC Lattice target groups will be routed to your Auto Scaling group.
Create new VPC Lattice service
Application Recovery Controller (ARC) zonal shift - new Info
During an Availability Zone impairment, target instance launches towards other healthy Availability Zones.
Enable zonal shift
New instance launches will be retargeted towards healthy Availability Zones until the zonal shift is canceled.
Health checks
Health checks increase availability by replacing unhealthy instances. When you use multiple health checks, all are evaluated, and if at least one fails, instance replacement occurs.
EC2 health checks
Always enabled
Additional health check types - optional Info
Turn on Elastic Load Balancing health checks
Recommended
Elastic Load Balancing monitors whether instances are available to handle requests. When it reports an unhealthy instance, EC2 Auto Scaling can replace it on its next periodic check.
EC2 Auto Scaling will start to detect and act on health checks performed by Elastic Load Balancing. To avoid unexpected terminations, first verify the settings of these health checks in the Load Balancer console
Turn on VPC Lattice health checks
VPC Lattice can monitor whether instances are available to handle requests. If it considers a target as failed a health check, EC2 Auto Scaling replaces it after its next periodic check.
Turn on Amazon EBS health checks
EBS monitors whether an instance's root volume or attached volume stalls. When it reports an unhealthy volume, EC2 Auto Scaling can replace the instance on its next periodic health check.
Health check grace period Info
This time period delays the first health check until your instances finish initializing. It doesn't prevent an instance from terminating when placed into a non-running state.
300
seconds
Cancel
Skip to review
Previous
Next

Step-6:In configure group and scaling,desired capacity is selected as 2,along with min value as 2 and max as 3.Target tracking scaling policy is selected.Instance warmup is set for 200 sec.

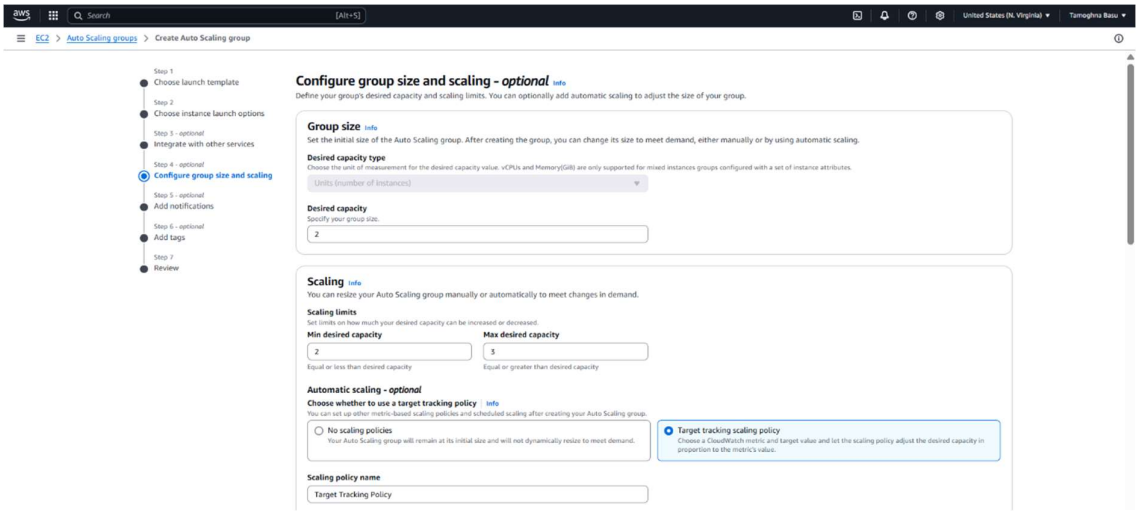

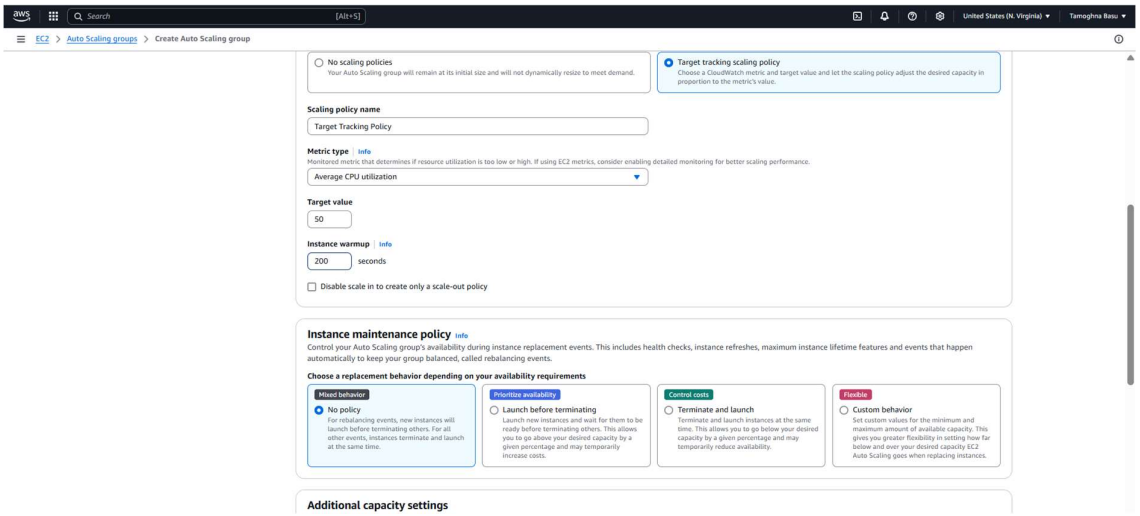

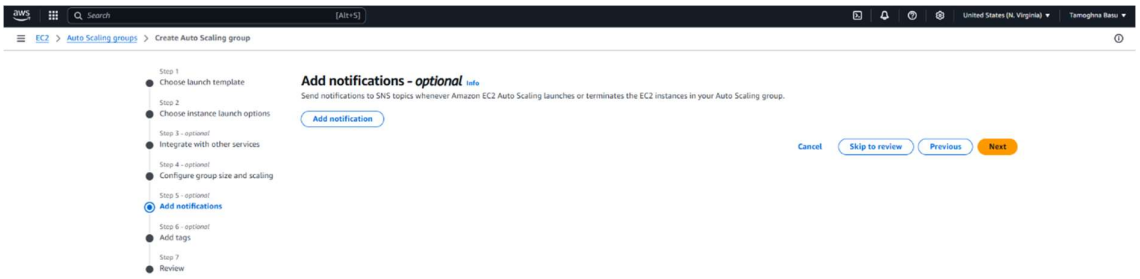

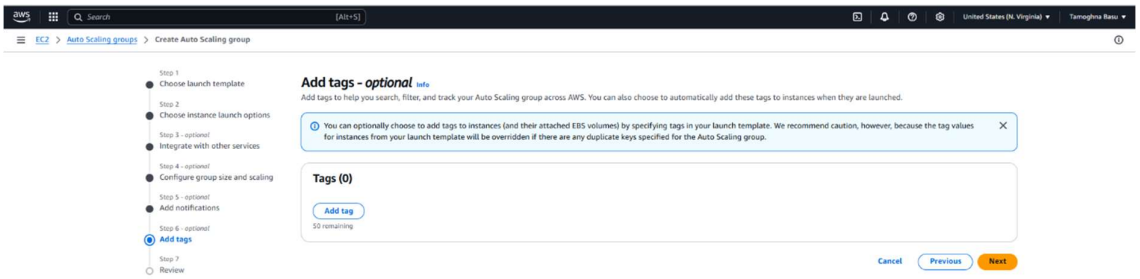

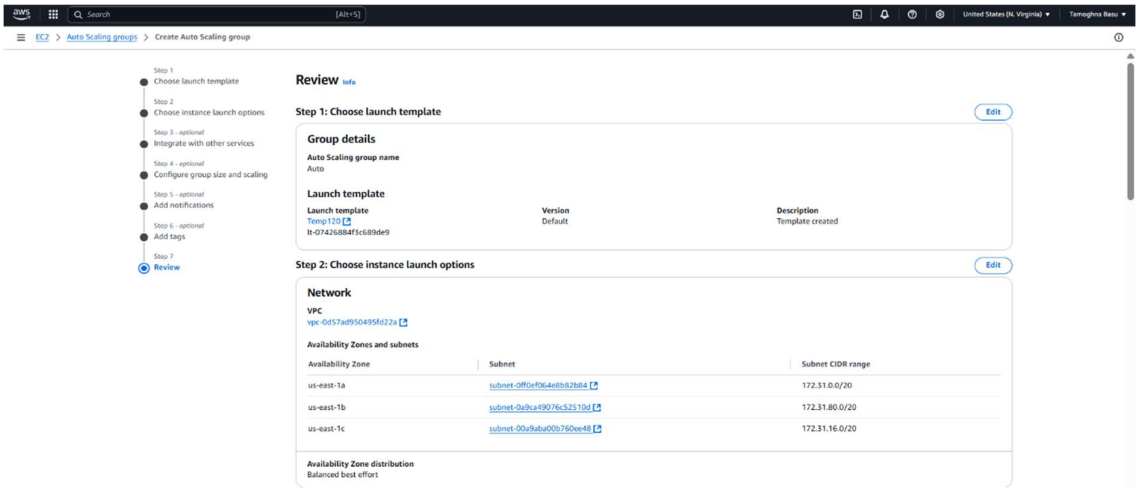

Step-7:Without making further modifications,simply clicking on next ,finally the auto scaling group has been created.

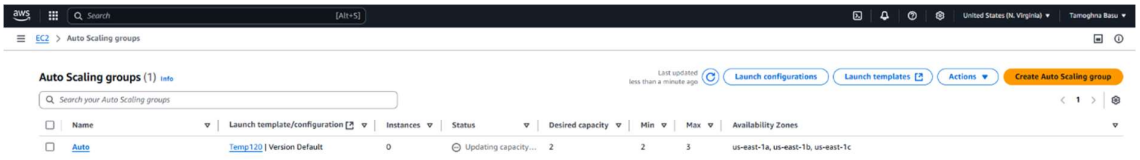

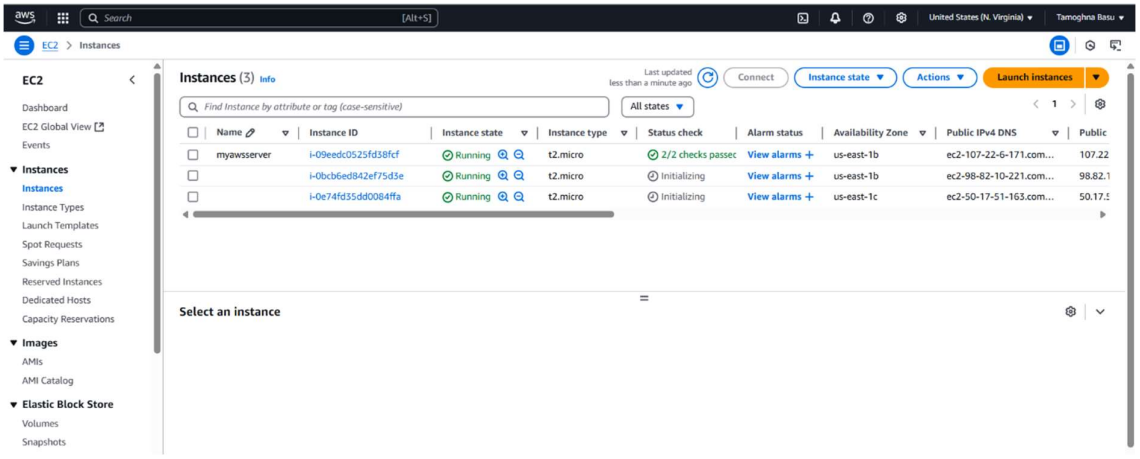

Step-8:Bitvise is opened,properly logged in with the key-pair and in the terminal following commands are written.

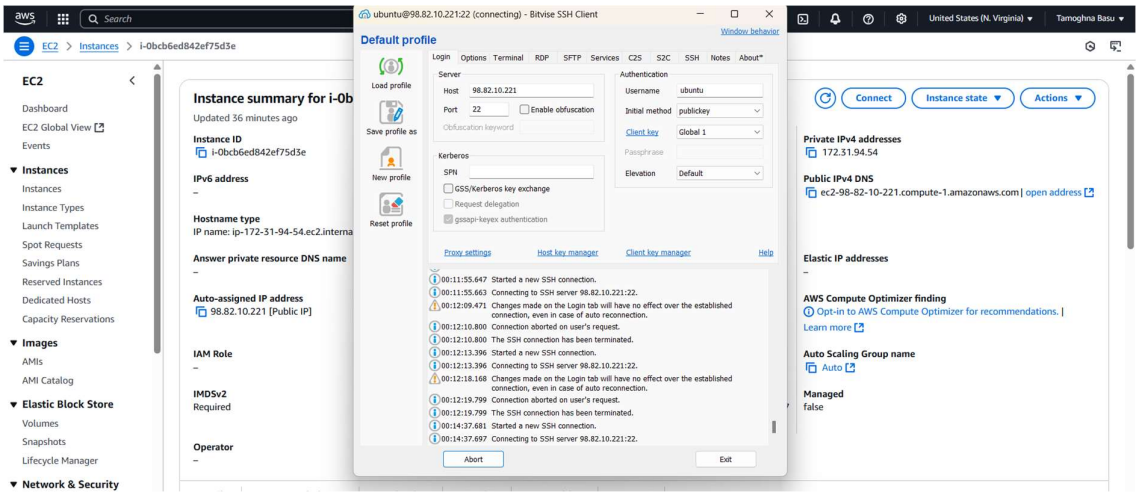

//Sudo nano ri.sh

Inside the file we have to write

```
#!/bin/bash
while(true)
do
    echo "Inside loop"
done
//sudo chmod +x ri.sh
//sh ri.sh
```

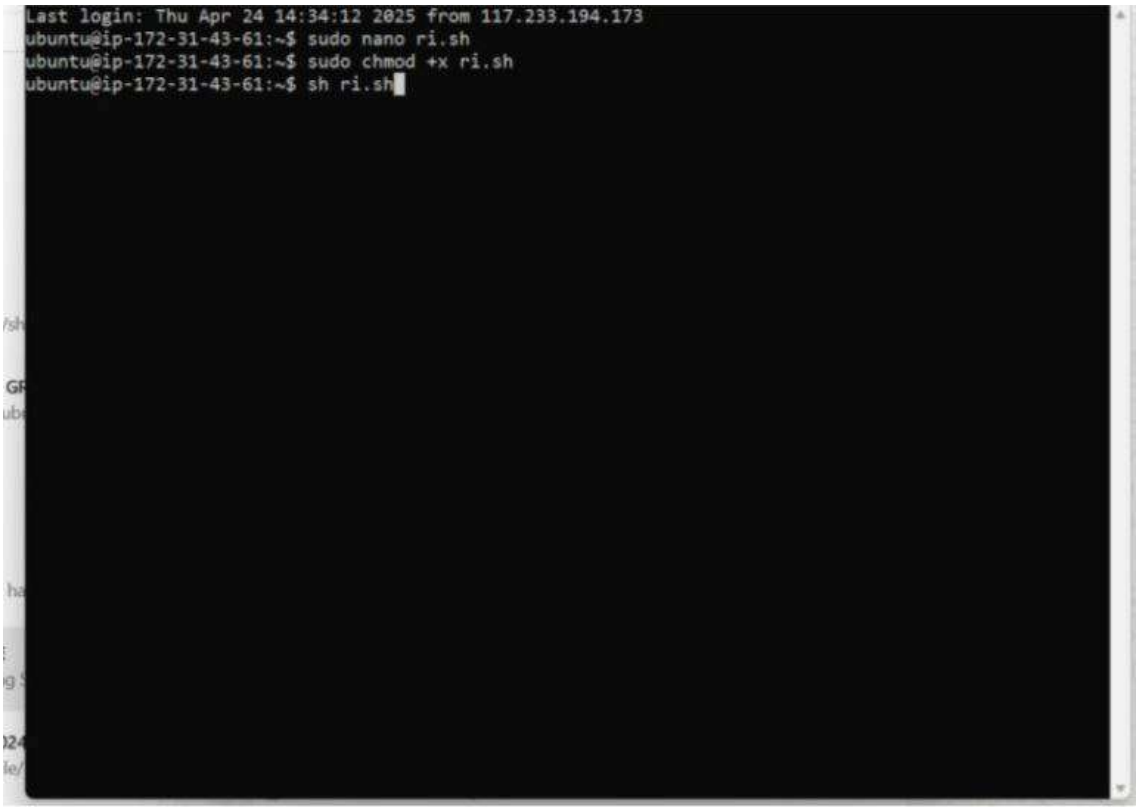

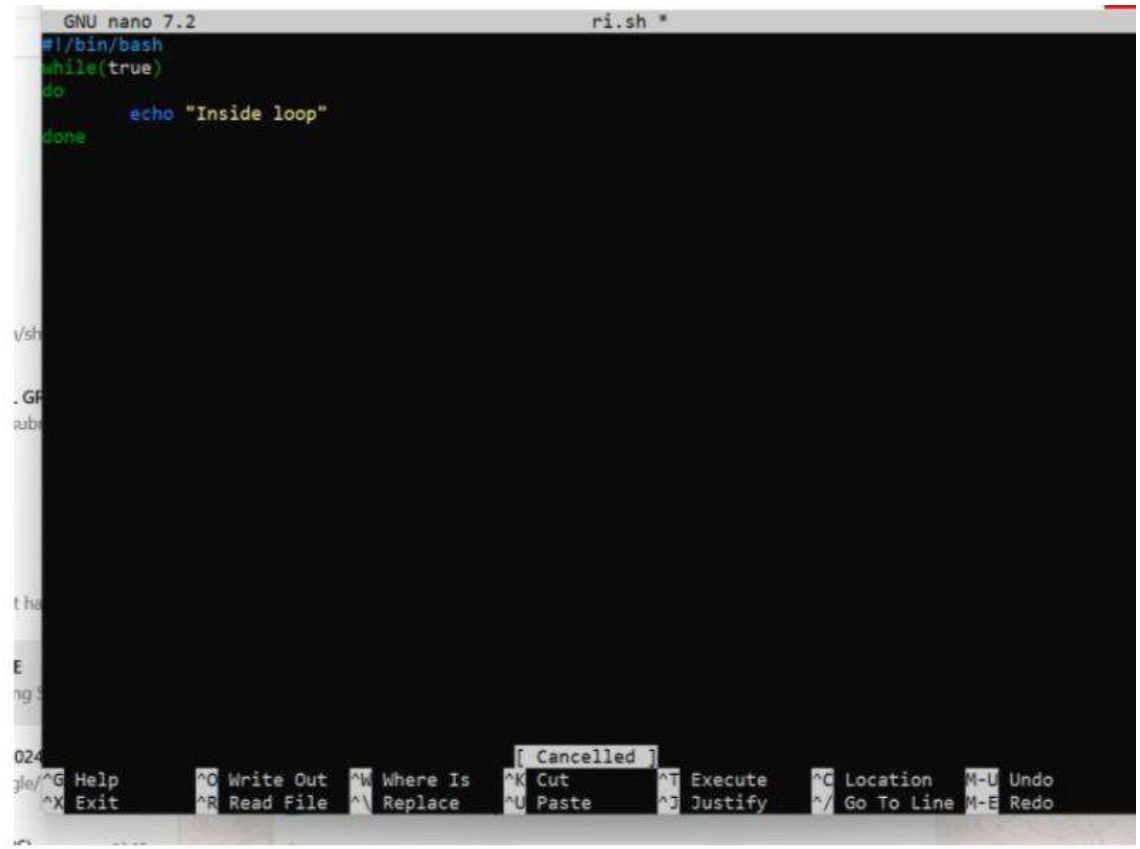

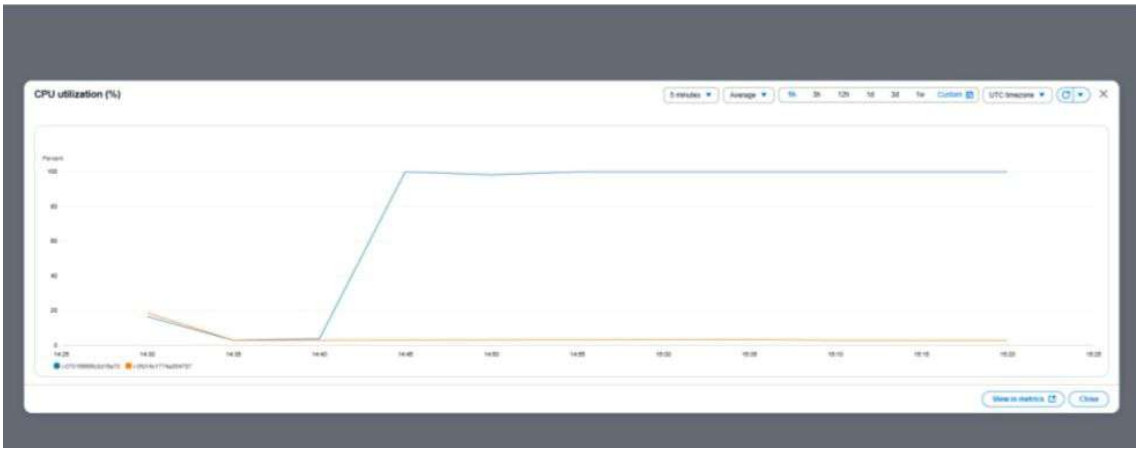

Step-8:CPU utilization is checked on the monitoring section after the bash file is executed and the third instance is created as a result of it.

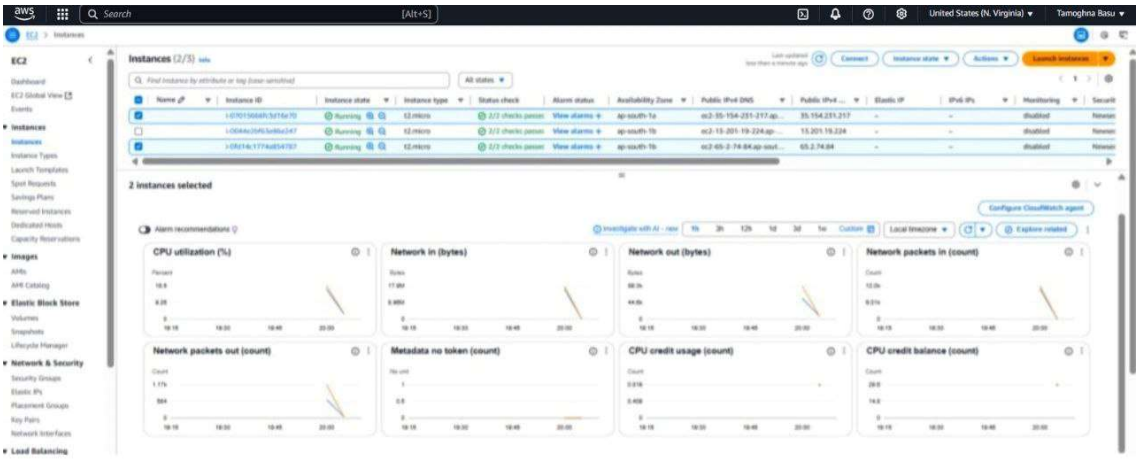
Search
[Alt+S]
United States (N. Virginia)
Tamoghna Basu
EC2
Instances
Instances (2/5)
Launch instances
2 instances selected
CPU utilization (%)
Network in (bytes)
Network out (bytes)
Network packets in (count)
Network packets out (count)
Metadata no token (count)
CPU credit usage (count)
CPU credit balance (count)